

AI Tool Kart — News Agent

Parameter, Threshold and Gate Summary

Quick reference. Companion to AGENT_ARTICLE_JUDGEMENT_CRITERIA.pdf (55 pp).


All values read from the implementation; identical to the detailed document.

Decision-source key: **D** = deterministic code • **L** = decided by the LLM • **L→D** = LLM-generated, code-enforced
Gate key: **block** = prevents publication • **defer** = pauses, story survives • **advisory** = reported only

Contents

1 Pipeline Overview	1
2 Discovery / Source Parameters	2
3 Freshness / Initial Pre-filter	2
4 Deduplication	3
5 Ranking / Candidate Selection	3
5.1 Scoring components	3
5.2 Equations (actual implemented weights)	3
5.3 Selection gates and caps	4
6 Source Trust and Evidence	4
7 Claim Extraction	4
8 Claim Verification	5
8.1 Support levels	5
8.2 Evidence requirement matrix (CLAIM_RULES)	5
8.3 Story-level gates (all must pass)	5
9 Article Format Selection	6
10 SEO Parameters	6
11 Writer Parameters	7
12 Editorial Validator	7
13 Category and Tag Parameters	8
13.1 Categories — fixed allowlist (D)	8
13.2 Tags — open vocabulary, validated (L→D)	8
14 WordPress Publication Gate	8
15 Cost and Automation Parameters	9
16 Final Approval Matrix	10

1 Pipeline Overview



Stage	Type	Decision made
Discovery	D	Which feeds are read; per-source item cap
Pre-filter	D	Freshness, banned domain, title length, hard-rejected topic, topic prior
Deduplication	D	Canonical-URL identity; title-similarity clustering
Ranking	mixed L→D	LLM scores relevance/importance; code computes trust + freshness, applies weights and floors
Evidence	D	Which pages are fetched, in what order; sufficiency rule
Claim extraction	L	Atomic claims + claim type (bounded by closed schema)
Verification	mixed L→D	LLM assigns support level; tier rules downgrade; code applies 4 story gates
Format selection	D	brief / standard / analysis from verified-claim and publisher counts
SEO brief	mixed L→D	LLM proposes metadata; code re-derives slug, filters links, validates grounding
Writer	L	Prose only, from verified claims; code renders HTML, slug, tags
Editorial validation	mixed L→D	LLM verdict; code overrides on deterministic issues and confidence floor
Taxonomy	D	Category allowlist; tag validation, normalisation, cap
WordPress draft	D	9 final checks; status = draft hardcoded

2 Discovery / Source Parameters

Parameter	Current value / rule	Src	Blocking?	Failure behaviour
Source registry	17 registered, 10 enabled (7 Tier 1, 3 Tier 2); single file <code>src/sources/registry.ts</code>	D	—	Adding a source is config, not code
enabled	true false per source	D	block	Source is not read at all
id	Stable string, e.g. <code>openai-news</code>	D	—	Used for tier/publisher lookup
trustTier	1 2 3; travels with every item into verification	D	block	Drives evidence + claim rules
Trust score by tier	Tier 1 = 10, Tier 2 = 7, Tier 3 = 3	D	—	Feeds the weighted score
type	rss api web ; only rss implemented & enabled	D	block	All web entries disabled
maxItemsPerRun	Per-source override; set to 15 for GitHub Changelog and AWS ML Blog	D	defer	Excess items not ingested this run
AGENT_MAX_ITEMS_PER_SOURCE_PER_RUN	Default 25 (global fallback)	D	defer	Excess items not ingested this run
publishedAt	From feed; used for freshness	D	—	See §3 for missing/invalid dates
Source fetch failure	Isolated per source; counted in sourcesFailed	D	no	Run continues with other sources
Source parse failure	Same isolation; body never logged	D	no	Run continues
Fetch order	Sequential, not parallel	D	—	—
HTTP guard	HTTPS only, private/link-local addresses rejected, 3 redirects max, AGENT_HTTP_TIMEOUT_MS = 15000, AGENT_MAX_RESPONSE_BYTES = 2500000	D	block	URL rejected; URL_REJECTED

Disabled sources (reason recorded, not guessed URLs): Anthropic, Meta AI, Mistral, Microsoft, Runway (Tier 1 — no stable feed / 404 / 410); Ars Technica (Tier 2 — feed valid, section too broad). **No Tier 3 source is configured**; Tier 3 remains the default tier for any unrecognised host.

3 Freshness / Initial Pre-filter

All checks below run **before the first LLM call** and are fully **D**.

Check / config name	Value / rule	Gate	Reason code on failure
maxStoryAgeHours	96 h , measured from the <i> freshest </i> backing item	block	stale (<i>nh</i> old)
Missing / unparseable date	publishedAt omitted; age falls back to discoveredAt (treated as fresh)	—	—
Future-dated item	Rejected beyond 6 h clock skew; date dropped, falls back to discoveredAt	—	—
minTitleLength	12 characters; checked at ingestion and again at prefilter	block	title-too-short
BANNED_DOMAINS	9 hosts (<code>news.google.com</code> , <code>medium.com</code> , <code>msn.com</code> , <code>yahoo.com</code> , <code>flipboard.com</code> , <code>benzinga.com</code> , <code>investing.com</code> , <code>analyticsinsight.net</code> , <code>marktechpost.com</code>)	block	banned-domain
EXCLUDED_TOPICS (hard)	5 rules: stock , doom , clickbait , listicle , horoscope-noise	block	excluded-topic: id
EXCLUDED_TOPICS (soft)	3 rules: rumor −3, celebrity −3, unrelated-funding −1.5	advisory	Score penalty only
INCLUDED_TOPICS	10 rules, weights +1 to +2.5 (highest: mcp +2.5)	advisory	Score bonus only
AUDIENCE_ENTITIES	43-entry vocabulary; +0.5 per hit, capped at 3 hits (+1.5 max)	advisory	Score bonus only
Terminal story status	Not published , not rejected , no duplicateOfStoryId	block	already-published , previously-rejected , duplicate-of-existing-story
Has backing items	≥ 1 news item	block	no-source-items

4 Deduplication

Level / parameter	Rule as implemented	Src
Canonical URL	Forced https ; host lowercased with www./m./amp. stripped; ~40 tracking params removed (utm_* , fbclid , gclid , pk_* , ...); redirect wrappers unwrapped one level ; AMP/ index.html suffixes stripped; remaining params sorted; fragment discarded	D
L1/L2 uniqueness	Canonical URL is a database uniqueness constraint on news_items	D
Title normalisation	Publisher masthead stripped; stop-words removed; 30+ announcement verbs folded to one token (launch/update)	D
Title similarity	$\max(0.7 \cdot \text{token Jaccard} + 0.3 \cdot \text{bigram Jaccard}, 0.85 \cdot \text{containment})$	D
sameStoryThreshold	0.62 — at or above this, one story	D
ambiguousThreshold	0.50 — between 0.50 and 0.62, merge but flag ambiguousMerge	D
Entity promotion	≥ 2 shared distinctive tokens promotes a near-miss (score ≥ 0.50) to a match	D
clusterWindowHours	72 h — titles only cluster within this window	D
Story fingerprint	Sorted normalised title tokens $\rightarrow$ stable story id across runs	D
Same-publisher guard	An <i>ambiguous</i> -grade match is refused when every existing item shares the new item's registrable domain	D

What happens to duplicates: an item whose canonical URL is already stored is dropped with no further work; an item whose title matches an existing story is *attached to that story* rather than starting a new one, so one event yields at most one article. Bias is toward merging (a wrong merge loses one article; a wrong split publishes a duplicate). Published and rejected stories still absorb duplicates, so a re-report cannot become a second article.

5 Ranking / Candidate Selection

5.1 Scoring components

Component	Range	Src	How produced
Relevance R	0–10	L	Classifier: would a reader act differently?
Importance I	0–10	L	Classifier: significance of the development
Novelty	0–10	L	Produced and schema-validated, consumed by nothing
Freshness F	0–10	D	Linear decay from age of freshest item
Source trust T	0–10	D	Best tier + bounded corroboration
Topic prior P	unbounded	D	Included + excluded rule weights + entity bonus
Category	enum(10)	L→D	Classifier picks; unrecognised $\rightarrow$ industry
Recommendation	proceed/skip	L	Advisory only ; logged, never decisive

5.2 Equations (actual implemented weights)

$$S_{\text{base}} = 0.35 R + 0.30 I + 0.20 T + 0.15 F \quad (1)$$

$$F(a) = \text{clamp}_{[0,10]} \left(10 \left(1 - \frac{a}{96} \right) \right), \quad a = \text{hours since freshest item} \quad (2)$$

$$T(\mathcal{T}, p) = \text{clamp}_{[0,10]} \left(\text{Base}(\min \mathcal{T}) + \min \left(\max(p - 1, 0) \cdot 1, 2 \right) \right), \quad \text{Base} = \{10, 7, 3\} \quad (3)$$

$$P = \sum_{\text{included}} w_r + \sum_{\text{excluded}} w_r + 0.5 \cdot \min(|\text{entity hits}|, 3) \quad (4)$$

$$A = \text{clamp}_{[-2,+2]} (0.4 P) \quad \Rightarrow \quad S = \text{clamp}_{[0,10]} (S_{\text{base}} + A) \quad (5)$$

p = independent publishers, counted by registrable domain. Corroboration bonus is **+1 each, capped at +2**, so many rewrites cannot outrank one primary source.

5.3 Selection gates and caps

Parameter	Value	Src	Purpose	Result on failure
MIN_RELEVANCE	4	D	Checked <i>first</i> ; trust/freshness must not carry an irrelevant story	block below-relevance-floor
MIN_IMPORTANCE	3	D	Filters trivial version bumps	block below-importance-floor
AGENT_MIN_WEIGHTED_SCORE	6.0	D	Bar for spending evidence budget	block below-threshold
Source base	Tier 1, or ≥ 2 independent publishers	D	An uncorroborated signal is a rumour	block insufficient-source-base
AGENT_MAX_CANDIDATES_PER_RUN	20	D	Bounds classifier spend	defer candidate-budget-reached
Topic-prior clamp	± 2	D	Keyword rules may shade, never overturn	—
Ordering	Descending S	D	Best material gets the remaining budget	—

All four selection gates are **conjunctive**. Deferred stories are re-queued *ahead of* new candidates on the next run, but are re-checked against the 96 h staleness gate — a deferral that aged out is then rejected, not carried indefinitely.

6 Source Trust and Evidence

Tier	Meaning	Operational treatment	Score
1	Vendor, changelog, official docs	Sufficient alone for evidence. Authoritative for its own product’s pricing, capabilities, dates, benchmarks. Not authoritative about competitors.	10
2	Reputable technology press	Reliable that something happened, not for exact specifications. Two distinct Tier 2 publishers required where no Tier 1 exists.	7
3	Social, forums, aggregators	Never sufficient . A claim whose supporting URLs are all Tier 3 is forced to unsupported . No Tier 3 source is configured today.	3

Parameter / rule	Value	Src	On failure
MAX_EVIDENCE_FETCHES_PER_STORY	4 pages per story; no links followed out of fetched pages	D	Remaining items not fetched
Fetch order	Tier 1 first , then by recency	D	Budget spends on primary material
MAX_EVIDENCE_CHARS	24,000 chars retained per source	D	Truncated
Minimum usable text	200 chars of extracted article text	D	Counted as a failed fetch
Sufficiency rule	≥ 1 Tier 1 retrieved, or ≥ 2 distinct Tier 2 publishers	D	block no-evidence-retrieved / insufficient-evidence (no Tier 1, n independent Tier 2)
Primary-source preference	Tier 1 fetched first; Tier 1 required for pricing / capability / date / benchmark / quote claims	D	Claim downgraded (§8)
Single-source allowance	Permitted in prose with attribution ; does <i>not</i> count toward article length	L→D	—
Headline-only evidence	Never permitted . Feed summaries are classifier context only and never reach the writer	D	Story rejected
HTTP 403 / 404 / 410	Permanent; not retried . Contributes no evidence	D	Rejected if sufficiency lost
HTTP 5xx / 408 / 429 / time-out	Transient; 3 attempts, backoff 400 ms base → 4 s ceiling	D	Then counted as failed
Prompt-injection markers	Whole document quarantined , not filtered; persisted for review, never sent to a model	D	Reason string appends (n source(s) quarantined for prompt injection)

7 Claim Extraction

One LLM call **per evidence document** (never across documents), so each claim’s origin is unambiguous.

Parameter	Value / rule	Src
Claim types (closed enum)	launch, capability, pricing, date, benchmark, quote, funding, other — 8, a ninth cannot be invented	L→D
Max claims per source	25 (ExtractionSchema)	D
Claim text length	10–400 characters	D
supportingQuote	Optional, verbatim span, ≤ 400 chars	L
containsInstructions	Required boolean — source flagged if it tries to instruct a model	L
Evidence linkage	Every claim carries the source URL it came from; required	D
Support level at extraction	Always unsupported ; only the verifier may raise it	D
Duplicate claims	Identical claims across sources are merged into one claim with multiple URLs (this is the corroboration signal)	D
Instruction-shaped claims	Dropped deterministically before verification	D
Malformed / off-schema output	Strict schema; 3 attempts (one repair, one clean retry), then error	D
One source fails extraction	Isolated — costs that source’s claims, not the story	D
Zero claims extracted	block no-claims-extracted	D

No confidence field exists on a claim. Confidence is expressed only as the support level assigned at verification.

8 Claim Verification

8.1 Support levels

Status	Definition	Effect on the article
verified	≥ 2 independent sources, or one Tier 1 stating it directly	May be stated as fact; counts toward the claim floor and toward format length
single-source	Exactly one non-Tier-1 source, uncorroborated	May appear with attribution ; does <i>not</i> buy length
unsupported	No supplied evidence states it (includes plausible inferences)	Removed — the writer never sees it
conflicting	Sources make incompatible statements; conflictNote required	Reported as a disagreement naming both values and publishers, or omitted

8.2 Evidence requirement matrix (CLAIM_RULES)

Claim type	Requirement for verified	Single-source OK?	Downgrade / rejection rule
launch	Tier 1, or 2 independent Tier 2	yes, attributed	One Tier 2 only → single-source
capability	Tier 1 required	yes, attributed	No Tier 1 → single-source
pricing	Tier 1 required	yes, attributed	No Tier 1 → single-source
date	Tier 1 required	yes, attributed	No Tier 1 → single-source
benchmark	Tier 1 required	yes, attributed	No Tier 1 → single-source
quote	Tier 1 required ; else 1 Tier 2	yes, attributed	No Tier 1 → single-source
funding	Tier 1, or 1 independent Tier 2	yes, attributed	All Tier 3 → unsupported
other	Tier 1, or 1 independent Tier 2	yes, attributed	All Tier 3 → unsupported

Universal: all-Tier-3 support → **unsupported**. No usable supporting URL → **unsupported**. A claim the verifier ignored → **unsupported**. Supporting URLs are filtered against the story’s real evidence set, so an echoed URL cannot manufacture support.

Tier rules only ever downgrade. The deterministic layer (**D**) runs *after* the LLM verdict (**L**) and can lower a support level but never raise one. Individual claim support is **advisory** at claim level (it shapes wording); the four story-level gates below are **block**.

8.3 Story-level gates (all must pass)

Gate	Src	Reason code on failure
Core claim (nominated by the verifier as coreClaimId) is not unsupported	L→D	core-claim-unsupported
Core claim is not conflicting without a Tier 1 resolution	L→D	core-claim-conflicting
≥ 3 usable claims (MIN_VERIFIED_CLAIMS)	D	too-few-verified-claims ($n < 3$)
≥ 1 fully verified claim	D	no-fully-verified-claims

Other bounds: MAX_CLAIMS_SENT = 25 claims per verification call; evidence excerpt 6,000 chars per source; response schema caps verdicts at 30.

Note: a **core: true|false** flag exists per claim type in config but **no code reads it** — the core claim is whichever claim the verifier nominates.

9 Article Format Selection

Fully **D**. Computed **once** before writing; the writer, the editor and any revision all use the same value.

Format	Words	Sections	Evidence required (all conjunctive)	SEO headings	SEO 2 nd kw
brief	180–350	3–5	The floor — no threshold to clear	3	2
standard	500–900	4–7	≥6 verified claims, ≥2 sources, ≥2 independent publishers	6	4
analysis	900–1400	5–9	≥12 verified claims, ≥3 sources, ≥3 publishers, importance ≥ 7	8	6

Signal	Role
verified claim count	Counts toward depth
single-source claim count	Does not count — weak sourcing cannot buy length
Evidence sources retrieved	Counts
Independent publishers	Counts , deduplicated by publisher name
Classifier importance	Counts for analysis only (≥ 7)
Internal links per format	2 / 3 / 4 (brief / standard / analysis)
ARTICLE.hardMaxWords	1400 — a draft above this is rejected as runaway output (block)
Below the format range	advisory below-format-target — expected when evidence runs out; padding is not a fix
Above the format range	advisory above-format-target — the more suspicious signal

10 SEO Parameters

Runs *after* verification and *before* writing. Fields are **L**-proposed, **D**-enforced.

Field	Constraint	Src
primaryKeyword	2–80 chars; must be grounded in verified claims	L→D
secondaryKeywords	Schema max 6 ; trimmed to format shape (2 / 4 / 6)	L→D
searchIntent	informational commercial navigational mixed	L
seoTitle	10–80 chars (schema); soft target 60	L→D
metaDescription	50–160 chars, enforced by schema	L→D
suggestedSlug	Proposal only — code re-derives with <code>slugify()</code> , max 72 chars	D
suggestedHeadings	Schema max 8 ; trimmed to format shape (3 / 6 / 8)	L→D
internalLinkTargets	Max 4 ; chosen from a 5-route allowlist (/browse, /workflows, /compare, /new-launches, /blog) plus this agent’s own published slugs. Unknown targets are silently dropped	D

BLOCKING SEO CHECKS — all **D**, evaluated against the verified claim set

seo-ungrounded-keyword	Primary keyword not supported by verified claims
seo-misleading-title	SEO title not supported by verified claims
seo-unsupported-superlative	Any of 24 ranking words (<i>best</i> , <i>top</i> , <i>fastest</i> , <i>leading</i> , <i>#1</i> , <i>revolutionary</i> , <i>industry-leading</i> , ...) in the SEO title, meta description or a keyword, unless the same word appears in a verified claim
seo-keyword-stuffing	Primary keyword > 4 times across title + headings + meta (a <i>count</i> , not a density); or a secondary keyword repeating the primary verbatim
seo-malformed-slug	Slug does not normalise to a usable URL
seo-missing-meta-description	Required metadata absent
seo-missing-primary-keyword	Required metadata absent

ADVISORY SEO CHECKS — reported, never blocking

seo-meta-description-short / -long	Outside 50–160 chars
seo-title-long	Over 60 chars
seo-keyword-not-in-title / -heading / -intro	Keyword placement (heading check skipped for briefs)
seo-heading-not-grounded	Suggested heading has no verified material
seo-too-many-headings	Beyond the format shape
seo-too-many-secondary-keywords	Beyond the format shape
seo-intent-mismatch	Commercial intent declared for a brief

Monotonic rule: SEO may *downgrade* an approved article to **needs-revision**, but can never upgrade a factually rejected article. Formally: $V' = V$ if $V \neq \text{approved}$; otherwise **needs-revision** when blocking SEO findings exist, else **approved**.

SEO failure is never fatal: a failed or skipped brief (e.g. no verified claims to ground it) degrades to *no brief* and the writer proceeds.

11 Writer Parameters

Parameter	Value / rule	Src
Model class	strong (TASK_MODEL_CLASS.write)	D
Max output tokens	4096 (TASK_MAX_OUTPUT_TOKENS.write)	D
Temperature	0.3	D
Inputs supplied	verified + single-source + conflicting claims (with type, support level, publishers, conflict note); source metadata; assigned format and word range; category; SEO brief if present; editor notes on a revision	D
Inputs withheld	Raw source documents, fetched page text, feed summaries, quarantined sources, and all unsupported claims	D
Structured output	Closed strict schema: title, excerpt, category, tags, sections, usedClaimIds . No HTML, no Markdown, no slug	D
Title	10–110 chars; must be supported by the claims	L→D
Excerpt	80–320 chars, written deliberately (not truncated from the body)	L→D
Tags	Schema 1–8; normalised and capped at 5 by code	L→D
Sections	3–9 ; 5 prescribed headings, omit rather than pad	L→D
usedClaimIds	Max 30; code discards any id not matching a real claim	D
Banned hype language	15 phrases (<i>in the rapidly evolving landscape, game-changing, revolutionary, delve into, unlock the power, a testament to, in conclusion, ...</i>); forbidden in the prompt and blocked by code	L→D
HTML rendering	Built by code from validated sections; 11-tag allowlist (p, h2, h3, ul, ol, li, a, strong, em, blockquote, code); all text escaped	D
Slug	Deterministic from title; fixed on first generation, reused on every revision	D

12 Editorial Validator

Check	Src	Gate	Rule / issue kind
Factual grounding	L	block	Every factual sentence must map to a claim — unsupported-claim
Exaggeration	L	block	“faster” → “3× faster” — overstated-claim
Headline accuracy	L	block	Body must establish the headline — misleading-headline
Attribution	L	block	Single-source claims attributed; conflicts reported — missing-attribution
Duplicate risk	L	block	Compared against last 100 published titles — duplicate
Tone / hype	L	advisory	hype — usually minor
Structure / grammar	L	advisory	structure, grammar — usually minor
Banned phrases	D	block	15 phrases — banned-phrase:"p"
Sources section	D	block	missing-sources-section, no-source-urls
Claim traceability	D	block	≥ 1 real claim id — no-claim-traceability
Excerpt	D	block	≥ 80 chars — excerpt-too-short
Tags (max)	D	block	> 5 — too-many-tags
Tags (min)	D	advisory	< 2 — too-few-tags (degrades gracefully in the UI)
HTML allowlist	D	block	disallowed-html-tags:tags
Article format / length	D	advisory	below-format-target / above-format-target
Category	D	block	Must be in the 10-term allowlist
SEO blocking findings	D	block	Folded in by the monotonic rule
Confidence	L→D	block	MIN_APPROVAL_CONFIDENCE = 0.60 — low-confidence

Decision state	Condition and consequence
approved	No model blocking issue, no deterministic blocking issue, no blocking SEO finding, confidence ≥ 0.60. Only this verdict may reach WordPress.
needs-revision	Blocking issues a rewrite could fix. One revision may be attempted.
rejected	Not fixable from the available claims.

Revision parameter	Value
MAX_REVISION_ATTEMPTS	1 (max 2 editorial reviews per article)
Model class / output tokens	strong / 2048 ; temperature 0
Revision triggered only when	verdict is needs-revision and ≥ 1 blocking issue is writer-fixable and the allowance is unused
A revision may not change	the assigned format, the SEO brief, or the slug
Editorial call fails entirely	Treated as rejected — editorial-review-unavailable (never fails open)
Code overrides the model	An approved verdict is lowered to needs-revision by any deterministic blocking issue or confidence < 0.60
Not approved	Draft is persisted with its verdict, confidence and full issue list ; story marked generated with editorial:needs-revision or editorial:rejected . Nothing is sent to the CMS

13 Category and Tag Parameters

13.1 Categories — fixed allowlist (D)

The 10 terms: ai-models, ai-agents, development, design, image-video, productivity, research, mcp, product-updates, industry.

Rule	Behaviour
Allowlist check (before any network call)	A category outside the 10 cannot reach WordPress. Classifier and writer both choose from a closed enum; unrecognised → industry
Category missing in WordPress	defer category-not-configured — the article waits; the term is never created
Arbitrary category creation	Forbidden. Seeding is an explicit operator action (<code>taxonomy:bootstrap</code>)
Per post	Exactly one category; resolved by slug and cached for the run

13.2 Tags — open vocabulary, validated (L→D)

Parameter	Value / rule	Gate
Min / max tags	2 (advisory) / 5 (blocking)	advisory / block
Length	2–40 chars; too-short , too-long	reject tag
Word count	Max 3 words; too-many-words	reject tag
Character class	Letters, digits, spaces and . + # ' & / -; disallowed-characters	reject tag
Control chars / markup	Checked before tidying; control-characters , markup	reject tag
Theme words	60+ blocked terms (“ai”, “tools”, “news”, “launch”, “best”, months, headline verbs); theme-not-entity	reject tag
Slug	Must produce a usable slug; unslugifiable	reject tag
Canonical vocabulary	43 entries; “open ai”/“openai”/“OpenAI” → <i>OpenAI</i> . Exact lookups only, no fuzzy matching	—
Must appear in the article	Every accepted tag (curated hits included) must be named in the draft, word-boundary matched	reject tag
Unlisted vendors	Accepted if proper-noun shaped, short, not a theme word, and present in the text	—
Deduplication	By slug (“OpenAI” and “openai” are one term)	—
Tag creation	WORDPRESS_CREATE_TERMS (default true); needs only edit_posts	—
Tag creation fails (incl. 403)	Tag is skipped , never fatal ; the article publishes with whichever tags resolved	advisory

14 WordPress Publication Gate

All checks **D**. Evaluated in order before any post is created.

#	Requirement	On failure
1	Editorial status approved	block skipped — not-approved (<i>status</i>)
2	No blocking SEO issue	Already folded into the verdict; cannot be bypassed
3	Category on allowlist and present in WordPress	defer category-not-configured / taxonomy-resolution-failed
4	Rendered HTML within the 11-tag allowlist	block at editorial stage
5	Sources block present + ≥ 1 source URL	block missing-sources-section, no-source-urls
6	No existing <code>wp_post_id</code>	Skipped — already-published (<code>wp:n</code>)
7	Story is not a duplicate of a published story	Skipped — duplicate-of-published-story
8	Run draft budget available	defer article stays pending for the next run
9	WordPress configured, reachable, authenticated	defer wordpress-not-configured / wordpress-unavailable. 401/403 is not retried and disables publishing for the whole run

`status = draft`

The agent does not auto-publish. The status is hardcoded in the payload builder (not read from config); `AGENT_AUTO_PUBLISH` is validated **false** at startup and the agent **refuses to start** if set true; the HTTP client asserts the status again and refuses any non-draft post. Publication to readers is a human action.

Idempotency: the article row is re-read before the request; merged-duplicate stories are aborted; the returned post id is written to the database immediately and the repository refuses to overwrite it.

Pending retry: runs *first*, before ingestion. `Batch = min(PENDING_PUBLISH.maxPerRun = 10, AGENT_MAX_DRAFTS_PER_RUN)`; aborts after **2** consecutive failures; costs **zero LLM calls** and never regenerates the article.

15 Cost and Automation Parameters

Environment / config variable	Default	Meaning
AGENT_ENABLED	true	Kill switch. false exits cleanly <i>before</i> the database, provider or network are touched
AGENT_AUTO_PUBLISH	false	Must stay false — true makes startup fail
AGENT_MAX_ITEMS_PER_SOURCE_PER_RUN	25	Items taken from any one feed
AGENT_MAX_CANDIDATES_PER_RUN	20	Stories sent to the classifier
AGENT_MAX_STORIES_VERIFIED_PER_RUN	5	Stories taken through evidence + verification
AGENT_MAX_ARTICLES_PER_RUN	2	Articles written
AGENT_MAX_DRAFTS_PER_RUN	2	Total CMS posts per run, retries included
AGENT_MAX_LLM_CALLS_PER_RUN	60	Absolute circuit breaker on calls
AGENT_MAX_TOKENS_PER_RUN	250000	Absolute circuit breaker on tokens
AGENT_MIN_WEIGHTED_SCORE	6	Selection threshold (0–10)
AGENT_RUN_LOCK_STALE_MINUTES	30	Age at which a run lock is reclaimed
AGENT_HTTP_TIMEOUT_MS	15000	Per-request timeout
AGENT_MAX_RESPONSE_BYTES	2500000	Response size cap
WORDPRESS_CREATE_TERMS	true	Governs tags only; categories never created
AGENT_DB_PATH	./data/news-agent.sqlite	Dedupe, idempotency, audit trail
AGENT_USER_AGENT	AIToolKartNewsAgent/0.1	Identifies the agent to publishers
AGENT_LOG_FORMAT / AGENT_LOG_LEVEL	pretty / info	json for production
LLM_PROVIDER	mock	mock runs offline and deterministically; any other value requires a key or startup fails
LLM_MODEL_FAST / LLM_MODEL_STRONG	adapter default	Per-class model overrides
<i>Credential variables LLM_API_KEY, WORDPRESS_API_URL, WORDPRESS_USERNAME, WORDPRESS_APP_PASSWORD exist and are required for live runs. Listed by name only — values are never printed, never logged (registered for redaction at startup) and never exposed to the browser bundle.</i>		
<i>Per-task output ceilings (TASK_MAX_OUTPUT_TOKENS) and model class (TASK_MODEL_CLASS)</i>		
classify / extract / seo	512 / 2048 / 1024	model class fast
verify / write / edit	3072 / 4096 / 2048	model class strong

Semantic	Behaviour
Cap unset / blank	Documented default applies
Cap = 0	Valid: zero work. The run defers remaining work and exits cleanly
Cap invalid (negative, fractional, non-numeric)	Startup fails loudly , naming the variable. Caps never silently fall back to a default — a cost control must fail closed
Deferred on budget	Story keeps its place and is re-queued <i>ahead of</i> new candidates next run, subject to a fresh staleness re-check
Deferral reason codes	candidate-budget-reached, verification-budget-reached, article-budget-reached, llm-budget-exhausted
What caps never stop	Publishing an article an earlier run already approved (costs no LLM budget). Set AGENT_MAX_DRAFTS_PER_RUN=0 to stop all CMS writes
Run lock	One invocation = one run, no daemon. A concurrent run exits 0 having done nothing; a lock older than 30 min is reclaimed

16 Final Approval Matrix

Stage	Key parameter(s)	Pass condition	On failure
Automation	AGENT_ENABLED	true	Run exits 0, no work done
Concurrency	Run lock, 30 min stale	No active run	Run exits 0
Source	enabled, trustTier, type	Source enabled and parseable	Source skipped; run continues
Domain	BANNED_DOMAINS (9)	Host not banned	block banned-domain
Title	minTitleLength	≥ 12 chars	block title-too-short
Freshness	maxStoryAgeHours	Freshest item ≤ 96 h	block stale (<i>nh</i> old)
Topic	5 hard-reject rules	No hard-reject match	block excluded-topic: <i>id</i>
Dedupe (URL)	Canonical URL uniqueness	Not already stored	Item dropped as duplicate
Dedupe (title)	0.62 / 0.50, 72 h window	No matching recent story	Merged into that story
Ranking	$R \geq 4$, $I \geq 3$, $S \geq 6$, Tier 1 or ≥ 2 publishers	All four	block below-relevance-floor, below-importance-floor, below-threshold, insufficient-source-base
Candidate cap	20 per run	Under cap	defer candidate-budget-reached
Evidence	4 fetches, 200-char floor, tier rule	≥ 1 Tier 1 or ≥ 2 Tier 2 publishers retrieved	block no-evidence-retrieved / insufficient-evidence
Integrity	Injection markers	No source attempts injection	Source quarantined; rejected if sufficiency lost
Extraction	8 claim types, max 25/source	≥ 1 claim extracted	block no-claims-extracted
Verification	CLAIM_RULES, MIN_VERIFIED_CLAIMS=3	Core claim supported; ≥ 3 usable; ≥ 1 verified	block core-claim-unsupported, core-claim-conflicting, too-few-verified-claims, no-fully-verified-claims
Format	6/2/2 and 12/3/3/imp 7	Highest fully-met tier assigned	Falls back to brief (not a failure)
Word ceiling	hardMaxWords = 1400	Draft ≤ 1400 words	block rejected as runaway output
Writer	Verified claims only, closed schema	Schema-valid draft	block after 3 attempts
Editorial	Model verdict + 7 deterministic checks	approved, no blocking issue	block retained — editorial:needs-revision / editorial:rejected
Confidence	MIN_APPROVAL_CONFIDENCE=0.60	≥ 0.60	block lowered to needs-revision
Revision	MAX_REVISION_ATTEMPTS=1	Writer-fixable, allowance unused	Retained for human review
SEO	7 blocking checks	No blocking finding	block approved verdict $\rightarrow$ needs-revision
Taxonomy	10-category allowlist; tags ≤ 5	Category valid and exists; tags validated	defer category-not-configured; block too-many-tags; individual tags skipped
Budget	Articles 2, drafts 2, calls 60, tokens 250k	Under every cap	defer story keeps its place
WordPress	9 final checks	All pass	Skipped or deferred; the draft is never lost
Output	status = draft	—	Human review required before publication

Design preference: *no article* is always preferred to *unsupported article*. Every ambiguous case resolves toward not publishing — unreadable evidence ends the story, unestablished claims are removed rather than softened, a review that cannot run is a rejection, and an exhausted budget pauses work rather than lowering the bar.